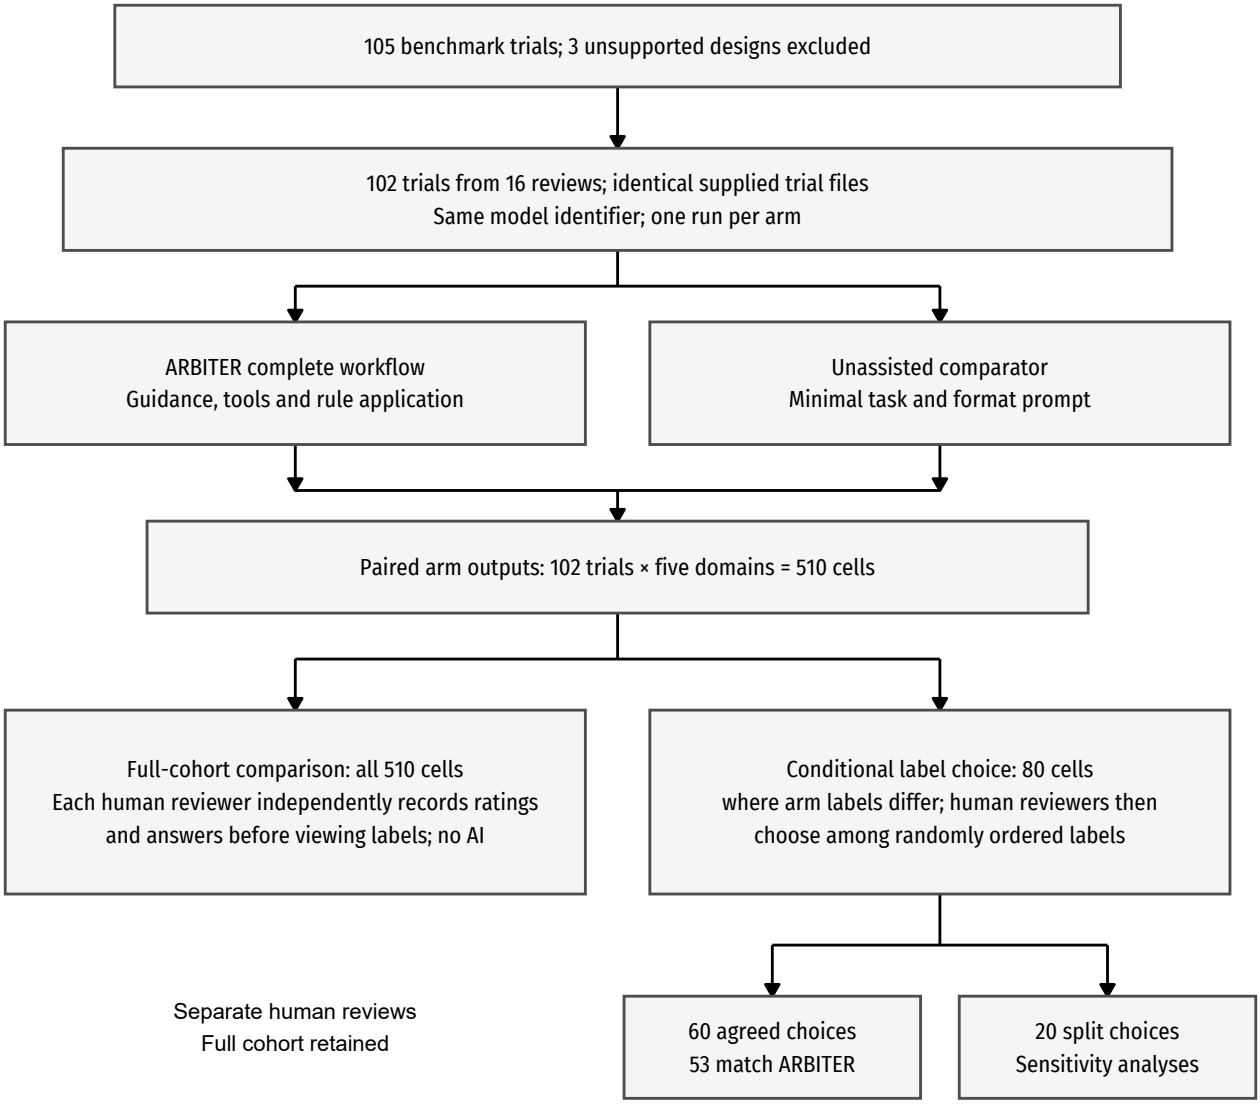
105 benchmark trials; 3 unsupported designs excluded
102 trials from 16 reviews; identical supplied trial files
Same model identifier; one run per arm
ARBITER complete workflow
Guidance, tools and rule application
Unassisted comparator
Minimal task and format prompt
Paired arm outputs: 102 trials × five domains = 510 cells
Full-cohort comparison: all 510 cells
Each human reviewer independently records ratings and answers before viewing labels; no AI
Conditional label choice: 80 cells
where arm labels differ; human reviewers then choose among randomly ordered labels
Separate human reviews
Full cohort retained
60 agreed choices
53 match ARBITER
20 split choices
Sensitivity analyses