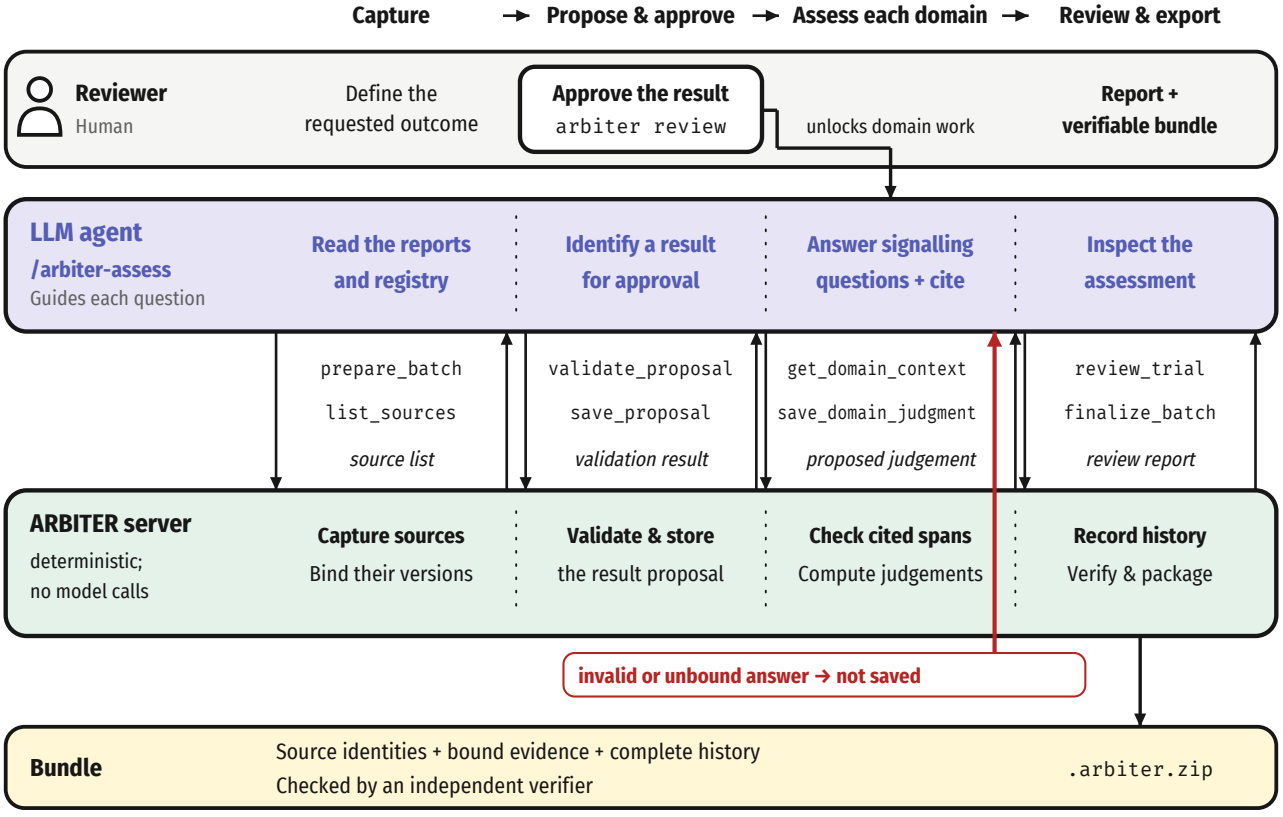
Capture
Propose & approve
Assess each domain
Review & export
Reviewer
Human
Define the requested outcome
Approve the result
arbiter review
unlocks domain work
Report + verifiable bundle
LLM agent
/arbiter-assess
Guides each question
Read the reports and registry
Identify a result for approval
Answer signalling questions + cite
Inspect the assessment
prepare_batch
list_sources
source list
validate_proposal
save_proposal
validation result
get_domain_context
save_domain_judgment
proposed judgement
review_trial
finalize_batch
review report
ARBITER server
deterministic;
no model calls
Capture sources
Bind their versions
Validate & store
the result proposal
Check cited spans
Compute judgements
Record history
Verify & package
invalid or unbound answer → not saved
Bundle
Source identities + bound evidence + complete history
Checked by an independent verifier
.arbiter.zip