Table 1. ARBITER detected more of the risk flagged by published reviews, at lower specificity

Source	Weighted κ	Exact agreement	Sensitivity[1]	Specificity
Against Published labels — D1–D5 (510 cells); Some concerns or High: N = 83; Low N = 427				
ARBITER	0.34 (0.26–0.42)	75.7% (72.4–79.0)	66.3% (57.3–74.7)	81.3% (77.7–84.7)
Unassisted	0.28 (0.21–0.35)	77.5% (74.1–80.8)	50.6% (40.3–60.2)	85.7% (82.1–89.3)
All-Low rule	—	83.7% (81.0–86.5)	—	100.0%
Against Published labels — Overall (81 trials); Some concerns or High: N = 39; Low N = 42				
ARBITER	0.44 (0.32–0.56)	58.0% (49.4–66.7)	92.3% (83.7–100.0)	54.8% (44.2–67.4)
Unassisted	0.42 (0.27–0.56)	59.3% (50.6–67.9)	64.1% (50.0–77.8)	76.2% (65.9–86.8)
All-Low rule	—	51.9% (44.4–59.3)	—	100.0%
Paired difference, ARBITER − Unassisted				
Published: D1–D5	0.06 (−0.03 to 0.14)	−1.8 (−5.1 to 1.6)	15.7 (6.8–24.4)	−4.4 (−7.6 to −1.2)
Published: Overall	0.02 (−0.11 to 0.16)	−1.2 (−11.1 to 8.6)	28.2 (16.1–40.5)	−21.4 (−32.7 to −9.8)

Published means published review label. Values are coefficients or percentages (95% CI); denominators are in block labels. [1] Sensitivity for Some concerns or High; specificity is the share of published Low ratings also rated Low. All-Low rule has no variance; only accuracy and specificity are reported. Paired differences are coefficients or percentage points; — indicates undefined or not applicable. AC2 and composite-reference results: eTables 6 and 29. Full paired differences: eTables 7 and 30. Overall judgements were not adjudicated. Published labels were frequently less severe than the adjudicators' agreed choice (Figure 4). CI, confidence interval; 10,000 percentile bootstrap replicates, trials resampled within meta-analysis; seed 20261018.