Table 2. Before seeing any label, both adjudicators' own ratings agreed with ARBITER far more than with the unassisted model or the published review

Comparison	Weighted κ	Raw agreement
Between adjudicators		
Independent ratings	0.59 (0.48–0.68)	114/154 (74.0%; 68.2–79.9)
Ratings chosen after seeing labels	0.69 (0.57–0.79)	129/154 (83.8%; 78.6–88.9)
Adjudicator 1 own rating compared with …		
Published	−0.08 (−0.17 to 0.01)	34/154 (22.1%; 15.8–28.5)
ARBITER	0.52 (0.41–0.62)	107/154 (69.5%; 62.7–76.0)
Unassisted	0.22 (0.13–0.30)	69/154 (44.8%; 38.1–51.6)
Adjudicator 2 own rating compared with …		
Published	−0.08 (−0.18 to 0.03)	42/154 (27.3%; 20.4–34.3)
ARBITER	0.67 (0.56–0.76)	120/154 (77.9%; 71.7–84.0)
Unassisted	0.31 (0.19–0.42)	76/154 (49.4%; 42.2–56.5)

All comparisons use the 154 domain ratings reviewed because at least one Sonnet arm differed from the published label. These comparisons cannot estimate agreement or published-label accuracy in the full benchmark. Independent ratings preceded blinded label presentation. Adjudicators then chose among the three offered labels. Linearly weighted Cohen's κ; AC2 is in eTable 17; raw agreement is n/N (x.x%; 95% CI a–b). CI, confidence interval; 10,000 percentile bootstrap replicates, trials resampled within meta-analysis; seed 20261018.