Table 3. Without ARBITER, 1 in 8 domain ratings broke the RoB 2 rules

Measure	Unassisted	ARBITER
Domain ratings (D1–D5)		
Rating contradicts its own answers	26/509 (5.1%; 3.3–6.9)	0 by design
Required question skipped	39/509 (7.7%; 5.7–10.0)	0 by design
Any rule break	65/509 (12.8%; 10.0–15.6)	0 by design
D1 rule breaks	5/101 (5.0%; 1.0–8.9)	0 by design
D2 rule breaks	7/102 (6.9%; 2.9–10.8)	0 by design
D3 rule breaks	5/102 (4.9%; 1.0–8.8)	0 by design
D4 rule breaks	35/102 (34.3%; 26.5–42.2)	0 by design
D5 rule breaks	13/102 (12.7%; 6.9–18.6)	0 by design
Overall rating		
Overall rating follows the RoB 2 rules (incl. permitted escalation)	102/102 (100.0%)	102/102 (100.0%)
Evidence and audit		
Answers citing ≥1 supporting passage or image	Not available	1538/1538 (100.0%)
Quoted text linked to its source	Not available	1407/1407 (100.0%)
Image regions linked to their page	Not available	14/14 (100.0%)
Audit file passes the independent verifier	Not available	102/102 (100.0%)

Values are n/N (x.x%; 95% CI a–b) where an interval is shown. A rule break means a rating contradicts its own answers or a required question was skipped. ARBITER computes proposed ratings from the official rules and requires complete answers; recorded overrides: 0/510. One ARBITER Overall rating used the permitted escalation; both arms followed the Overall rules in every trial. Full checks and skipped-question sensitivity: eTables 23 and 24. One Unassisted domain lacks a rating and answers and is excluded (509 evaluable). Evidence coverage counts required answers citing a passage or image. Text and image checks count distinct cited records. ARBITER extracts quoted text and renders cited images. The model writes image descriptions; these are not checked against the images. The independent verifier checks the audit file structure, not clinical correctness or completion. Unassisted evidence records are unavailable. CI, confidence interval; 10,000 percentile bootstrap replicates, trials resampled within meta-analysis; seed 20261018.